

# 1 What Standard SGD Does

Stochastic Gradient Descent fits model parameters  $\theta$  by taking small steps that decrease the calculated loss over the dataset. At each step:

- Pick a random batch of  $b$  examples.
- Compute the average gradient.
- Update  $\theta$  by moving in the negative gradient direction.

---

**Algorithm 1** Standard SGD

---

**Require:** Data  $\{x_i\}_{i=1}^N$ , loss  $L(\theta, x)$ , initial  $\theta_0$ , learning rates  $\{\eta_t\}$ , batch size  $b$ , steps  $T$ .

```
1: for  $t = 1$  to  $T$  do  
2:   Sample mini-batch  $B_t$  of size  $b$ .  
3:    $g_t \leftarrow \frac{1}{b} \sum_{x \in B_t} \nabla_{\theta} L(\theta_{t-1}, x)$ .  
4:    $\theta_t \leftarrow \theta_{t-1} - \eta_t g_t$ .  
5: end for
```

---

## 2 Making SGD Private

DP-SGD ensures that no particular sample is able to have an outsized impact on the gradient at any step of iteration [2].

This is accomplished by clipping each examples gradient, to limit the impact of each sample, and then adding noise to hide individual contribution differences:

1. Clipping each per-example gradient to norm at most  $C$ .
2. Adding Gaussian noise of standard deviation  $\sigma C$  to the averaged gradient.

---

**Algorithm 2** Differentially Private SGD

---

**Require:** Data  $\{x_i\}$ , loss  $L(\theta, x)$ , initial  $\theta_0$ ,  $\{\eta_t\}$ ,  $b$ ,  $C$ ,  $\sigma$ ,  $T$ .

```
1: for  $t = 1$  to  $T$  do  
2:   Sample  $B_t$  of size  $b$ .  
3:   for each  $x \in B_t$  do  
4:      $g_t(x) \leftarrow \nabla_{\theta} L(\theta_{t-1}, x)$ .  
5:      $\bar{g}_t(x) \leftarrow g_t(x) / \max(1, \|g_t(x)\|_2 / C)$ .  
6:   end for  
7:    $\bar{g}_t \leftarrow \frac{1}{b} \sum_{x \in B_t} \bar{g}_t(x)$ .  
8:    $\tilde{g}_t \leftarrow \bar{g}_t + \mathcal{N}(0, \sigma^2 C^2 I)$ .  
9:    $\theta_t \leftarrow \theta_{t-1} - \eta_t \tilde{g}_t$ .  
10: end for
```

---

## 3 Various values of $\sigma$

There are various methods by which we quantify the amount of noise added. The tighter the bounds we need, the better.

Let  $q = b/N$  be the sampled proportion and we target overall  $(\varepsilon, \delta)$ -DP.

### 3.1 1. Naïve Composition Theorem

Each iteration is viewed as  $(\varepsilon', \delta')$ -DP, and composing  $T$  of them yields  $(T\varepsilon', T\delta')$ -DP. To achieve  $(\varepsilon, \delta)$ :

$$\varepsilon' = \frac{\varepsilon}{T}, \quad \delta' = \frac{\delta}{qT}.$$

The Gaussian mechanism (with sensitivity  $C$ ) then requires

$$\sigma_{\text{naive}} = \frac{\sqrt{2 \ln(1.25/\delta')}}{\varepsilon'} = \frac{qT \sqrt{2 \ln(\frac{1.25 qT}{\delta})}}{\varepsilon}.$$

as per the works of Dwork [2].

### 3.2 2. Strong Composition Theorem

Advanced composition gives

$$(\tilde{\varepsilon}, T\delta' + \delta'')\text{-DP}, \quad \tilde{\varepsilon} = \varepsilon' \sqrt{2T \ln \frac{1}{\delta''}} + T\varepsilon' \frac{e^{\varepsilon'} - 1}{e^{\varepsilon'} + 1}.$$

For  $\varepsilon' \leq 1$ , this simplifies and leads to

$$\sigma_{\text{strong}} = O\left(\frac{q \sqrt{T \ln(1/\delta) \ln(T/\delta)}}{\varepsilon}\right).$$

once again taken from the works of Dwork [2].

### 3.3 Note: Naïve vs. Advanced

We compare the two noise scales by their ratio:

$$\frac{\sigma_{\text{strong}}}{\sigma_{\text{naive}}} = O\left(\frac{q \sqrt{T \ln(1/\delta) \ln(T/\delta)}/\varepsilon}{qT \sqrt{2 \ln(1.25 qT/\delta)}/\varepsilon}\right) = O\left(\sqrt{\frac{\ln(1/\delta) \ln(T/\delta)}{2T \ln(1.25 qT/\delta)}}\right).$$

For large  $T$ ,  $\ln(T/\delta) \approx \ln T$  and  $\ln(1.25 qT/\delta) \approx \ln T$ , so

$$\frac{\sigma_{\text{strong}}}{\sigma_{\text{naive}}} = O\left(\sqrt{\frac{\ln T}{T}}\right) = O(T^{-1/2}).$$

Thus, advanced composition requires asymptotically  $\sqrt{T}$  times less noise.

### 3.4 Moments Accountant Method

While Advanced Composition provides a good bound, better bounds exist for the case of DP SGD [1]. In particular, it has been shown that a method known as the Moments Accountant Method gives a noise variance bound of:

**Theorem 1** (Moments Accountant, [1, Thm. 1]). *If*

$$\sigma \geq c \frac{q \sqrt{T \ln(1/\delta)}}{\varepsilon},$$

*then DP-SGD satisfies  $(\varepsilon, \delta)$ -DP.*

This removes the extra  $\sqrt{\ln(T/\delta)}$  factor, making this a tighter bound on noise than the Advanced composition theorem.

## 4 Training Longer and Learning Rates

- Multiple epochs ( $ET$  steps) scale privacy loss by  $E$  for the naïve method, by  $\sqrt{E}$  using advanced composition, and only by a constant factor when using the moments accountant method.
- Changing  $\eta_t$  affects convergence but not the privacy analysis: only the number of noisy steps and noise level matter.

## References

- [1] Martín Abadi, Andy Chu, Ian Goodfellow, H. Brendan McMahan, Ilya Mironov, Kunal Talwar, and Li Zhang. Deep learning with differential privacy. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*, pages 308–318, 2016.
- [2] Cynthia Dwork and Aaron Roth. *The Algorithmic Foundations of Differential Privacy*, volume 9 of *Foundations and Trends in Theoretical Computer Science*. Now Publishers, 2014.